



Yasir Rehman

@ yasirrehman1

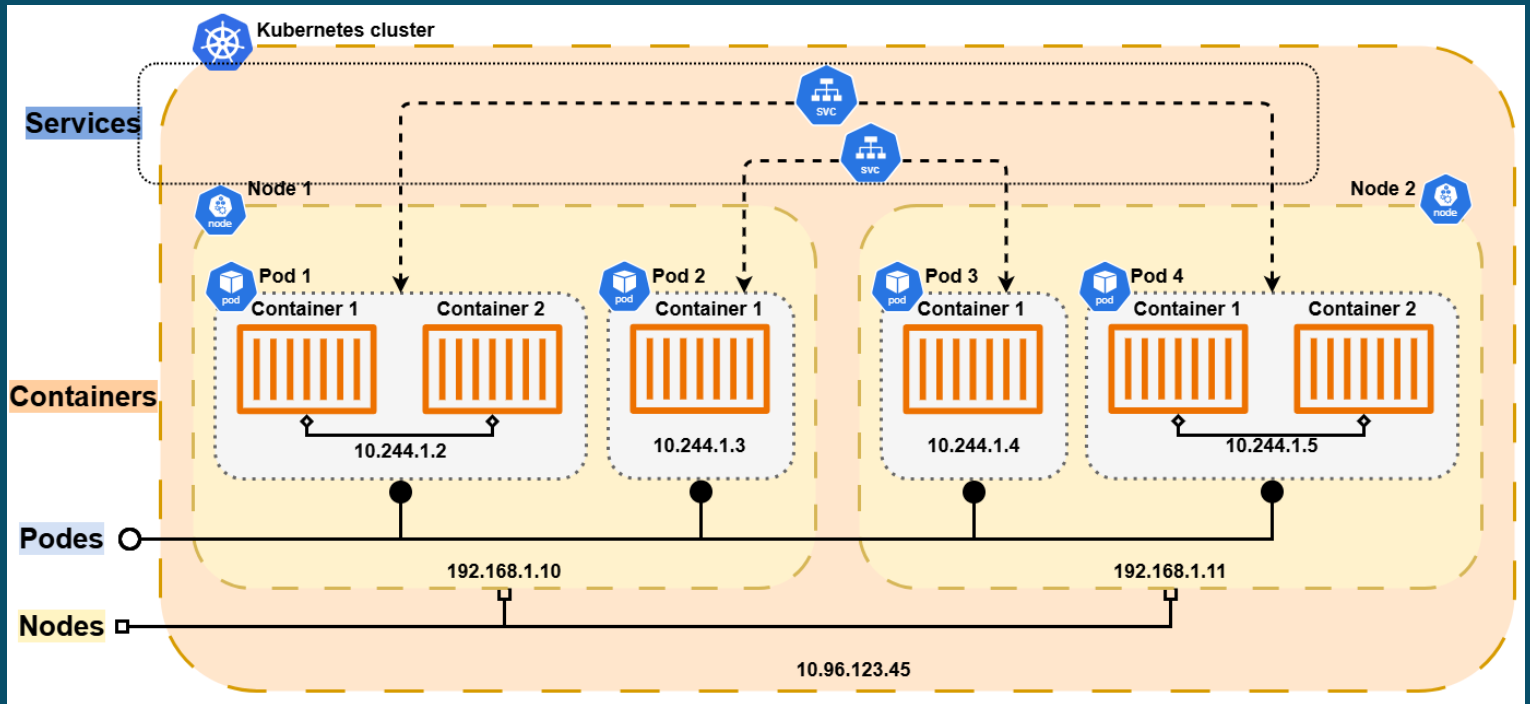


# Exploring Kubernetes Networking

TECHNOLOGY: KUBERNETES

LEVEL: INTERMEDIATE

Kubernetes **networking** is built on a few **core principles** that ensure simplicity, scalability, and resilience.



## Core principles:

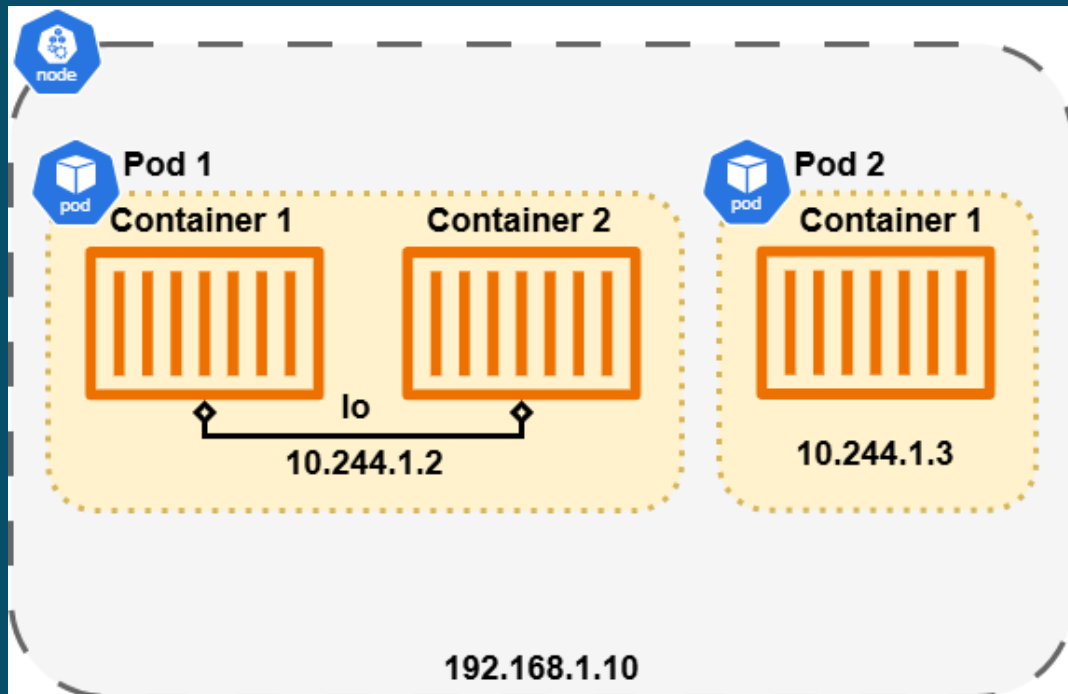
- Each **Pod** is assigned a **unique IP** address from the CIDR range of its network, **eliminating** the need for **NAT** and simplifying communication and debugging.
- Pods can **communicate** with each other **across nodes**, ensuring seamless operation and facilitating scalability as the cluster expands.

- **Services abstract Pods**, providing a **stable Cluster IP** that decouples applications, enables load balancing, and handles dynamic environments.
- Kubernetes' built-in **DNS** service allows Services to be discovered using **human-readable** DNS names, simplifying communication and eliminating the need for hardcoded IPs.
- Kubernetes offers several mechanisms—NodePort, LoadBalancer, and Ingress—to **expose applications externally**, making them accessible to users and systems outside the cluster.
- **CNI plugins**, such as Calico, Flannel, and Weave, manage Pod **IP assignment**, **routing**, and **network policies**, offering flexibility in choosing networking solutions tailored to specific needs.

# Networking scopes

## Container-level network

The container level network refers to the networking stack that is specific to an **individual container**.



Each container has its **own network namespace**, which includes:

- A **loopback interface** (lo) for internal communication.
- A **virtual Ethernet interface** (veth) that connects the container to the host's network stack.

## How it works?


- When a container is created, it is **assigned** a **virtual Ethernet interface** (veth), which is paired with a corresponding interface on the host.
- The container's **network namespace** is **isolated** from other containers, meaning it has its own IP address, routing table, and firewall rules.

## Key characteristics:

- Each container has its **own network namespace**.
- Containers within the **same Pod** share the **same network namespace**.

# Pod-level network

The Pod network is the network scope that enables **communication between Pods** in a Kubernetes cluster.



It is a **flat** network where:

- Every Pod gets a unique IP address.
- Pods can **communicate** with each other **directly**, regardless of which node they are running on.

## How it works?


- When a Pod is created, it is assigned an **IP** address from the **Pod network's CIDR** (Classless Inter-Domain Routing) range.
- The **CNI** (Container Network Interface) plugins such as Calico, Flannel, and Weave are responsible for managing **IP assignment**, **routing**, and **network policies** for Pods.

## Key characteristics:

- All Pods can communicate with each other **without NAT**.
- Pods on different nodes can **communicate seamlessly**.
- Each Pod has a **unique IP** address that is routable within the cluster.

# Node-level network

The node network refers to the **physical or virtual network** that connects the **nodes** (machines) in a Kubernetes cluster.



This network is responsible for:

- Communication between the **control plane** (e.g., API server, etcd) and **worker nodes**.
- Communication **between worker nodes** for Pod-to-Pod traffic across nodes.



## How it works?

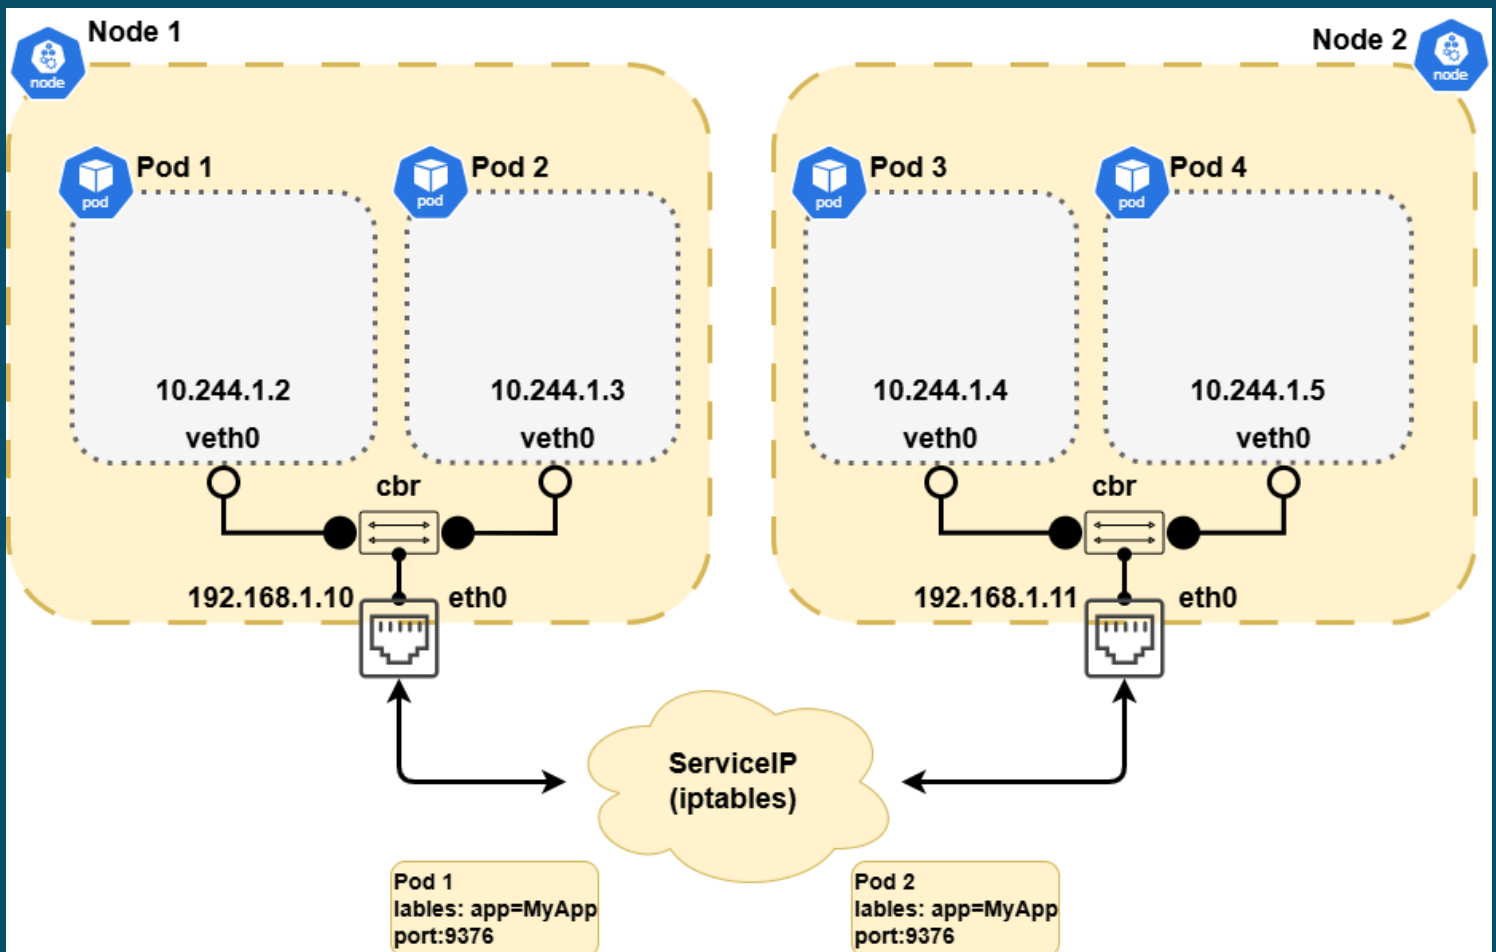
- Each node has a **physical or virtual network interface** (e.g., eth0) with an IP address assigned from the node network's CIDR range.
- The CNI plugin ensures that traffic between Pods on different nodes is routed correctly through the node network.

## Key characteristics:

- The node **network** is typically provided by the **underlying infrastructure** (e.g., cloud provider, on-premises network).
- Each node has a unique IP address that is used for communication between nodes.

# Service-level network

The service network is the layer in Kubernetes that **provides stable, abstracted endpoints** for accessing Pods.



Services act as an **intermediary between Pods and internal/external systems**, providing load balancing, service discovery, and external access.

## How it works?

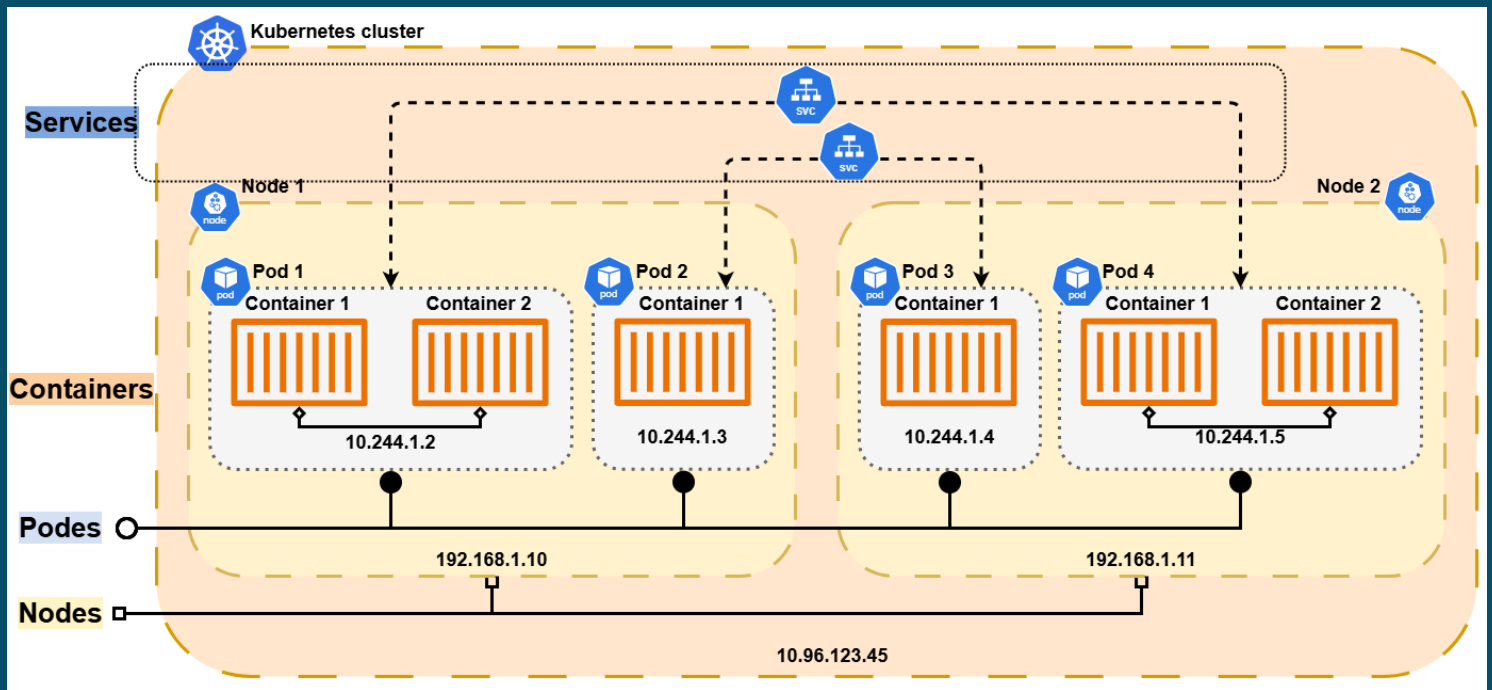
- A Service is **created with a selector** that matches the labels of the Pods it should route traffic to.
- The Service is **assigned a Cluster IP**, which is a virtual IP address that remains stable.
- The Service **dynamically updates its Endpoints** based on the Pods that match its selector.

## Key characteristics:

- Services are assigned **DNS** names in the **format** **<service-name>.<namespace>.svc.cluster.local**.
- kube-proxy uses **iptables or IPVS** to implement load balancing across the Pods.
- Traffic is **distributed** evenly across **backend Pods**, ensuring high availability and scalability.

# Cluster-level network

The cluster network is the overarching network that encompasses all the **networks** within a Kubernetes **cluster**.



It includes:

- The pod network (for **pod-to-pod** communication).
- The node network (for **node-to-node** communication).
- The service network (for **Service IPs**).

## How it works:

- The cluster network is **managed** by Kubernetes components like **kube-proxy** and the **CNI** plugin.
- kube-proxy ensures that **Services are reachable**, and that traffic is **routed** correctly to **Pods**.
- The **CNI** plugin manages **IP assignment**, **routing**, and **network policies** for Pods.


## Key characteristics:

- Services are assigned IPs from a separate CIDR range (e.g., 10.96.0.0/12).
- Kubernetes provides **DNS-based service discovery** within the cluster network.
- **Network policies** can be applied to **control traffic flow** within the cluster network.

# Communication scenarios

## Intra-Pod (container-to-container)

A **Pod** contains **multiple containers** that need to communicate with each other.




Containers within the **same Pod** share the **same network namespace** and can communicate using localhost(127.0.0.1).

**No additional** network configurations are needed for **intra-pod** communication.

```
apiVersion: v1
kind: Pod
metadata:
  name: multi-container-pod
spec:
  containers:
    - name: app
      image: nginx
    - name: sidecar
      image: busybox
      command: ["sleep", "3600"]
```

## Inter-Pod (pod-to-pod)

Pods on **different nodes** need to communicate with each other.



Pods are **assigned unique IP** addresses from the Pod network, and the **CNI** plugin **ensures** that **traffic** is **routed** correctly between nodes.


Communication happens **without Network Address Translation (NAT)**.

### Example: Direct pod communication

```
kubectl run pod1 --image=busybox --restart=Never --command -- sleep 3600
kubectl run pod2 --image=busybox --restart=Never --command -- sleep 3600
kubectl exec -it pod1 -- ping pod2-ip-address
```

## Pod-to-Service

A **Pod** needs to communicate with a **Service**.



The Service provides a stable Cluster IP, and **kube-proxy** ensures that **traffic is routed** to one of the backend Pods.

Since pod IPs are ephemeral, Services provide **stable endpoints**.


Example: Exposing a Deployment with a Service

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: my-app
  ports:
    - protocol: TCP
      port: 80
      targetPort: 8080
  type: ClusterIP
```



## External access (external-to-pod)

An application needs to be accessible from **outside the cluster**.




External traffic can reach the cluster through a **LoadBalancer or Ingress**, which routes the traffic to the appropriate Pods via the Service.

## Example: Ingress resource

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: my-ingress
spec:
  rules:
  - host: myapp.example.com
    http:
      paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: my-service
            port:
              number: 80
```

## DNS and service discovery

Services need to be discovered using DNS names.



Kubernetes provides **DNS-based service discovery**, allowing Pods to resolve Services using their DNS names.

Kubernetes provides **CoreDNS** for **internal** name resolution.

Services are accessible via ***service-name.namespace.svc.cluster.local***.

Example: Resolving a service name

```
kubect1 exec -it pod1 -- nslookup my-service.default.svc.cluster.local
```

# Network policies

Network policies allow you to **control traffic flow** between Pods.



Network policies **act as a firewall**, allowing or denying traffic based on labels and namespaces.

By default, **all pods can communicate** with each other.

Example: Denying all traffic except from a specific app

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: restrict-access
spec:
  podSelector:
    matchLabels:
      app: backend
  policyTypes:
    - Ingress
  ingress:
    - from:
        - podSelector:
            matchLabels:
              app: frontend
```



Yasir Rehman

@yasirrehman1

Find this insightful?

Repost! 

I am a Senior DevOps Engineer and Team Lead.

I post about DevOps, cloud-native technologies,  
and compassionate leadership.